

Lecture Notes for **Machine Learning in Python**



Professor Eric Larson
Dimensionality Reduction and Images

Class Overview, by topic

Table Data
Visualization

Numpy, Pandas, Seaborn
Overviews with some in-depth discussion

Dimension
Reduction and
Image Processing

Scikit-learn, Scikit Image,
Intuition only, Some mathematics

Linear and
Logistic
Regression

Numpy, Recreate API for Scikit-learn
Detailed mathematics for simple optimization
intuition for advanced optimization

Neural Networks
and Back Prop.

Numpy
Detailed mathematics for NN operations

Wide and Deep
Networks

Convolutional
Networks

Recurrent
Networks

Keras, Tensorflow
Intuition, Detailed implement.

Ethics in
Language Models

ConceptNet
Case studies

Class Logistics and Agenda

- **Logistics:**
 - **Next Time:** Flipped Module
 - Do **quiz one** after this lecture!!
 - Turn in one per team (HTML), please include team member names from canvas
- **Agenda**
 - Common Feature Extraction Methods for Images
 - Begin Town Hall, if time

Class Overview, by topic

Table Data
Visualization

Numpy, Pandas, Seaborn
Overviews with some in-depth discussion

Dimension
Reduction and
Image Processing

Scikit-learn, Scikit Image,
Intuition only, Some mathematics

Linear and
Logistic
Regression

Numpy, Recreate API for Scikit-learn
Detailed mathematics for simple optimization
intuition for advanced optimization

Neural Networks
and Back Prop.

Numpy
Detailed mathematics for NN operations

Wide and Deep
Networks

Convolutional
Networks

Recurrent
Networks

Keras, Tensorflow
Intuition, Detailed implement.

Ethics in
Language Models


ConceptNet
Case studies

Features of Images



Extracting Features: Convolution

- For images:
 - kernel (matrix of values)
 - slide kernel across image, pixel by pixel
 - multiply and accumulate



This Example:

3x3 Kernel (dark)

Ignoring edges of input

Input Image is 5x5

Output is then 3x3

Convolution

$$\sum \left(\mathbf{I} \left[i \pm \frac{r}{2}, j \pm \frac{c}{2} \right] \odot \mathbf{k} \right) = \mathbf{O}[i, j] \quad \text{output image at pixel } i, j$$

input image slice centered in i, j
with range $r \times c$

kernel of size, $r \times c$
usually $r=c$

0	0	0	0	0	0	0	0	0
0	1	2	3	4	12	9	8	0
0	5	2	3	4	12	9	8	0
0	5	2	1	4	10	9	8	0
0	7	2	1	4	12	7	8	0
0	7	2	1	4	14	9	8	0
0	5	2	3	4	12	7	8	0
0	5	2	1	4	12	9	8	0
0	0	0	0	0	0	0	0	0

input image, \mathbf{I}

0	0	0
2	3	4
2	3	4


1	2	1
2	4	2
1	2	1

kernel
filter, \mathbf{k}
3x3

20	21	36				
		
...
		
...
...
...
...
...

output image, \mathbf{O}

Convolution Examples




Blur

1	1	1
1	1	1
1	1	1




Vertical Edges

-1	0	1
-1	0	1
-1	0	1



Sharpen

0	-1	0
-1	5	-1
0	-1	0



Self test:

What does this do?

A. move left pixel to center



B. move right to center

C. blur

0	0	0
1	0	0
0	0	0


Common operations

- the gradient (2D derivative)






$$\|\nabla J\| = \sqrt{J_x^2 + J_y^2}$$

$$\theta = \tan^{-1}(J_x / J_y)$$






DAISY: same features, regardless of orientation



1. Select u, v pixel location in image and radius
2. Take histogram of average gradient magnitudes in circle
for each orientation $h_{\Sigma}(u, v)$
3. Select circles in a ring, R
4. For each circle on the ring, take another histogram $h_{\Sigma}(\mathbf{l}_O(u, v, R_1))$
5. Repeat for more rings
6. Save all histograms as "descriptors"
 $[h_{\Sigma}(\cdot), h_{\Sigma}(\mathbf{l}_1(\cdot, R_1)), h_{\Sigma}(\mathbf{l}_2(\cdot, R_1)) \dots]$
7. Can concatenate descriptors as "feature" vector at that pixel

Efficient DAISY, Orient x Circle Radius convolutions



take histogram of
convolved images
at points u, v

one
convolution
per
orientation

one convolve
per ring size

Daisy Operator at u_0, v_0 is
Concatenated ||Histograms||

take **normalized** histogram of magnitudes

$$\mathcal{D}(u_0, v_0) = \left[\begin{array}{l} \tilde{\mathbf{h}}_{\Sigma_1}^T(u_0, v_0), \\ \tilde{\mathbf{h}}_{\Sigma_1}^T(\mathbf{l}_1(u_0, v_0, R_1)), \dots, \tilde{\mathbf{h}}_{\Sigma_1}^T(\mathbf{l}_T(u_0, v_0, R_1)), \\ \tilde{\mathbf{h}}_{\Sigma_2}^T(\mathbf{l}_1(u_0, v_0, R_2)), \dots, \tilde{\mathbf{h}}_{\Sigma_2}^T(\mathbf{l}_T(u_0, v_0, R_2)), \end{array} \right]$$

$$\tilde{\mathbf{h}}_{\Sigma}(u, v) = \left\| \left[\mathbf{G}_1^{\Sigma}(u, v), \dots, \mathbf{G}_H^{\Sigma}(u, v) \right]^T \right\|$$

Tola et al. “Daisy: An efficient dense descriptor applied to wide- baseline stereo.” Pattern Analysis and Machine Intelligence, IEEE

An intuition for the future: DAISY workflow

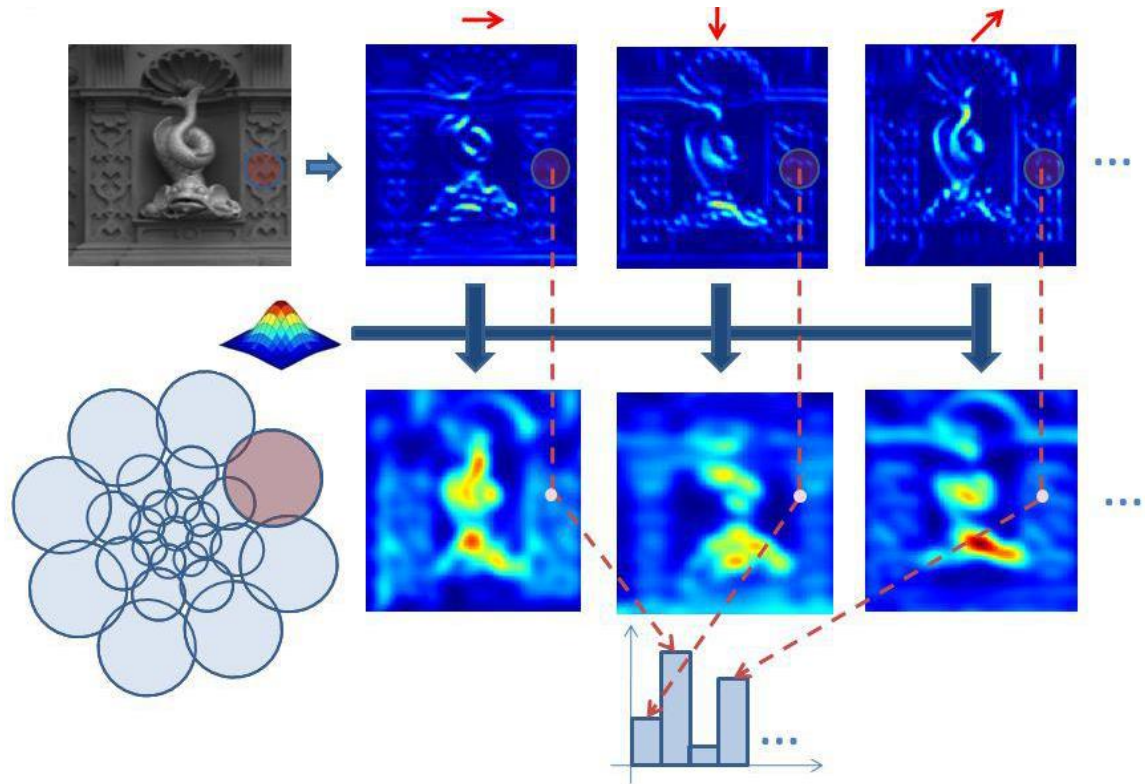
Convolutions on
Image



Convolutions of
Previous
Convolutions




Statistics
of Convolution
Magnitudes



Hyper Parameters in DAISY, need selection

Num circles/per ring



Num rings

Num orientations within each circle

Num of bins in histogram



```
daisy(img, step=180, radius=58, rings=2, histograms=6,  
      orientations=8, visualize=True)
```

Params

step, radius, num rings, num histograms per ring,
orientations, *bins per histogram*

Gradients

DAISY

(if time) Gabor Filter Banks




Other Tutorials:

http://scikit-image.org/docs/dev/auto_examples/

Matching versus Bag of Features

- Not a difference of vectors, but a percentage of matching points



- SURF, ORB, SIFT, DAISY

Feature Matching

Matching test image to source dataset

1. Choose src image from dataset
2. Take keypoints of src image
3. Take keypoints of test image
4. For each kp in src:
 1. Match with closest kp in test
2. *How to define match?*
5. Count number of matches between images
6. Determine if src and test are similar based on number of matches
7. Repeat for new src image in dataset
8. Once all images measured, choose best match as the target for the test image



match_descriptors

Scikit-image Implementation

```
skimage.feature.match_descriptors(descriptors1, descriptors2, metric=None, p=2,  
max_distance=inf, cross_check=True, max_ratio=1.0)
```

[\[source\]](#)

Brute-force matching of descriptors.

For each descriptor in the first set this matcher finds the closest descriptor in the second set (and vice-versa in the case of enabled cross-checking).